

A/B Testing MLOps Pipeline Implementation Plan

Overview

Transform your existing MLOps pipeline into a comprehensive A/B testing platform that can automatically experiment with different models, detect data drift, and retrain based on performance metrics from real user interactions.

Architecture Components

1. Traffic Splitting & Experiment Management

- **Traffic Router:** Intelligent routing between Model A (control) and Model B (treatment)
- **Experiment Manager:** Configure, start, stop, and analyze A/B experiments
- **Feature Flags:** Dynamic control of experiment parameters
- **Session Management:** Ensure consistent user experience during experiments

2. Data Drift Detection System

- **Drift Monitors:** Statistical tests for feature drift, prediction drift, and concept drift
- **Baseline Comparison:** Compare incoming data against training distribution
- **Alert System:** Automated alerts when drift exceeds thresholds
- **Drift Visualization:** Real-time dashboards showing drift metrics

3. Enhanced Monitoring & Metrics

- **A/B Metrics Collection:** Business metrics, model metrics, and user interaction data
- **Statistical Significance Testing:** Chi-square, t-tests, and confidence intervals
- **Real-time Dashboards:** Extended Grafana dashboards for A/B experiments
- **Performance Comparison:** Side-by-side model performance analysis

4. Automated Experiment Lifecycle

- **Experiment Design:** Automated experiment configuration based on business goals
- **Sample Size Calculation:** Statistical power analysis for experiment duration
- **Early Stopping:** Stop experiments early if significance is reached
- **Winner Selection:** Automated promotion of winning models to production

Infrastructure Layout

Existing Infrastructure (Reuse)

yaml

EKS Cluster: loan-eks-simple

Namespace: loan-default

MLflow: ab124afa4840a4f8298398f9c7fd7c7e-306571921.ap-south-1.elb.amazonaws.com

Grafana: (existing monitoring)

Prometheus: (existing alerts)

S3 Bucket: (existing DVC storage)

New Infrastructure Components

yaml

A/B Testing Namespace: ab-testing

Experiment Database: PostgreSQL (dedicated schema)

S3 Folders:

- experiments/data/
- experiments/models/
- experiments/results/

Traffic Splitter Service: Istio/NGINX-based routing

Drift Detection Service: Real-time monitoring

Implementation Phases

Phase 1: Core A/B Testing Framework (Week 1-2)

1. Traffic Splitting Service

- Implement intelligent traffic router
- Add experiment configuration management
- Create user session tracking

2. Experiment Database Schema

- Design experiments, metrics, and results tables
- Implement experiment lifecycle management
- Add audit trail and versioning

3. Basic A/B API Endpoints

- Experiment CRUD operations
- Traffic allocation endpoints
- Metrics collection APIs

Phase 2: Data Drift Detection (Week 2-3)

1. Drift Detection Engine

- Statistical drift tests (KS test, PSI, etc.)
- Feature distribution monitoring
- Prediction drift analysis

2. **Baseline Management**

- Training data distribution storage
- Dynamic baseline updates
- Historical drift tracking

3. **Alerting Integration**

- Prometheus metrics for drift
- Grafana dashboards for visualization
- Automated alert rules

Phase 3: Automated Experiment Management (Week 3-4)

1. **Experiment Automation**

- Automated experiment design
- Sample size calculations
- Statistical significance testing

2. **Model Comparison Pipeline**

- Automated A/B model training
- Performance comparison frameworks
- Winner selection algorithms

3. **Integration with Existing Pipeline**

- Trigger experiments from alerts
- Automated retraining based on A/B results
- MLflow experiment tracking enhancement

Phase 4: Advanced Analytics & Reporting (Week 4-5)

1. **Statistical Analysis Framework**

- Bayesian analysis for experiments
- Multi-armed bandit algorithms
- Causal inference tools

2. **Comprehensive Reporting**

- Experiment result reports
- Business impact analysis


- ROI calculations

3. Data Generation & Simulation

- Synthetic data generation with controlled drift
- Scenario testing capabilities
- Load testing with realistic data patterns

Technical Implementation Strategy

Data Flow Architecture



Database Schema Design

```
sql

-- Experiments table
experiments (id, name, start_date, end_date, status, config, created_by)

-- Experiment groups (A/B variants)
experiment_groups (id, experiment_id, name, model_version, traffic_allocation)

-- Metrics collection
experiment_metrics (id, experiment_id, group_id, metric_name, value, timestamp)

-- User assignments
user_assignments (user_id, experiment_id, group_id, assigned_at)

-- Drift measurements
drift_measurements (id, feature_name, drift_score, drift_type, measured_at)
```

Key Technologies Integration

- **MLflow**: Enhanced experiment tracking with A/B metadata
- **DVC**: Separate versioning for A/B models and datasets
- **Pytest**: Comprehensive testing for A/B logic and statistical functions
- **Prometheus**: A/B metrics and drift alerts
- **Grafana**: A/B dashboards and experiment monitoring

- **Kubernetes:** Container orchestration for A/B services



Data Drift Implementation

Drift Detection Methods

1. Feature Drift:

- Kolmogorov-Smirnov test for continuous features
- Chi-square test for categorical features
- Population Stability Index (PSI)

2. Prediction Drift:

- Model output distribution comparison
- Prediction confidence analysis
- Output stability metrics

3. Concept Drift:

- Performance degradation detection
- Label distribution changes
- Model accuracy decline patterns

Synthetic Data Generation

- **Controlled Drift Injection:** Gradually shift feature distributions
- **Scenario-based Testing:** Simulate different drift patterns
- **Realistic Data Patterns:** Maintain business logic consistency



Automated Workflows

A/B Testing Workflow

1. **Experiment Creation:** Define hypothesis, metrics, and success criteria
2. **Traffic Allocation:** Gradually ramp up traffic to treatment group
3. **Real-time Monitoring:** Track metrics and statistical significance
4. **Early Stopping:** Halt experiments if clear winner emerges
5. **Winner Promotion:** Deploy winning model to 100% traffic
6. **Post-experiment Analysis:** Generate comprehensive reports

Drift-based Retraining

1. **Continuous Monitoring:** Real-time drift detection
2. **Threshold Alerts:** Automated alerts when drift exceeds limits

3. **Experiment Trigger:** Launch A/B test with retrained model
4. **Performance Validation:** Compare new model against existing
5. **Automated Deployment:** Deploy if new model performs better

Testing Strategy

Unit Testing

- Statistical functions validation
- Traffic routing logic
- Drift detection algorithms
- Database operations

Integration Testing

- End-to-end A/B workflows
- MLflow integration
- Kubernetes deployment testing
- Alert system validation

Load Testing

- Traffic splitting under load
- Database performance with concurrent experiments
- Model serving performance comparison

Monitoring & Observability

Key Metrics Dashboard

- Experiment health and status
- Traffic distribution and balance
- Statistical significance tracking
- Business metrics comparison
- System performance metrics

Alerting Rules

- Experiment failures
- Statistical significance reached
- Drift threshold violations
- System performance issues

- Data quality problems

Expected Outcomes

Business Benefits

- Data-driven model improvement decisions
- Reduced risk of model deployments
- Faster innovation cycles
- Improved model performance
- Better understanding of user behavior

Technical Benefits

- Automated experiment lifecycle management
- Early detection of model degradation
- Reduced manual intervention
- Comprehensive audit trails
- Scalable experimentation platform

Next Steps

1. Review and approve the implementation plan
2. Set up the enhanced project structure
3. Implement Phase 1 components
4. Integrate with existing monitoring infrastructure
5. Launch first A/B experiment with synthetic data